

# Final Project Report for FA20 CSE208

Xiaohan Fu  
UC San Diego  
x5fu@ucsd.edu

## 1. Introduction

In this project, my major accomplishments are:

- Exploration on multiple FHE libraries, including SEAL, Lattigo, and Palisade;
- Better understanding of CMake and 3rd party libraries for C++;
- Implementation of an anonymous rider-driver pairing application with batching and serialization (code available [here](#)).

My initial plan was to use Lattigo as `golang` is more widely used in distributed systems. However after installed Lattigo and played with the samples a little bit, I found that the documentation is not that good and I felt a little bit confused with their APIs. Also, I felt I hadn't written C++ for a while so it might be a good idea to take a look at the C++ libraries. I studied Palisade, SEAL, and HELib and found their APIs to be generally similar. When building and installing Palisade, I met some problems that I couldn't resolve so I moved on to Microsoft SEAL. It has the best documentation (inline comments in source code) from my perspective as well. The installation process was non-trivial and interesting as I have never systematically used CMake and 3rd party libraries for C++ in the past. But fortunately everything worked well.

## 2. Implementation

The application simulates the situation where an anonymous rider needs to find the closest driver around to pick her up, hinted by the [sample code](#) of Lattigo [1]. On top of it, batching and serialization is utilized. The locations of the rider and drivers are represented as points in a rectangular grid. Server knows that the rider and some drivers are located within this area (this grid), but should not know their exact coordinates.

I strongly recommend to read the well-commented source code and run the application to see the output report for more details.

### 2.1. Workflow

The application is simulating the sense of server-client communications of real rider-driver pairing applications with two shared streams among users (drivers, rider) and servers for

data (ciphertext, public key) and parameter settings.

The workflow in general is:

1. The server publishes and sets the parameters of the FHE context;
2. The rider follows the parameters and generated a pair of **pk** and **sk**, then she uses her **sk** to symmetrically encrypt her coordinates and send **pk** and the ciphertext to the server;
3. The server broadcasts her **pk** to all the drivers in this region;
4. Drivers encrypt their own coordinates with the rider's **pk** and send them to the server;
5. The server evaluates the squared distance between the rider and each of the drivers homomorphically with relinearization after each operation and then send back **result** to the rider;
6. The rider decrypts **result** with her **sk**, and identifies the closest driver.

**Note.** Coordinates of the rider and drivers are randomly generated for each execution.

The security parameters chosen in Step 1 is:

polynomial modulus degree	ciphertext coefficient modulus	plaintext modulus
$N = 2^{13} = 8192$	218 (suggested by the helper function)	$T = 65929217$

A relatively small polynomial modulus was used here for small ciphertext size and faster execution. To support batching, the plaintext modulus  $T$  was chosen to be a prime number congruent to 1 modulo  $2 \times \text{poly\_modulus\_degree}$  ( $T = 65929217$  specifically here).

## 2.2. Batching

As hinted by the documentation of SEAL [2], I used batching for full utilization of the the plaintext polynomial. The total size of the batching vector is exactly the polynomial modulus degree  $N$ .

In Step 2, the plaintext that the rider is encrypting is not a single pair of her coordinates ( $\text{rider\_x}$ ,  $\text{rider\_y}$ ), but in the format of

$$[\text{rider}_x, \text{rider}_y, \text{rider}_x, \text{rider}_y, \dots, \text{rider}_x, \text{rider}_y].$$

Accordingly, in step 4, each driver is encrypting its coordinates in the format where for  $i$ -th driver, all coordinates are zero except  $(2i)$ -th and  $(2i+1)$ -th coordinates. For example, for  $\text{driver}[1]$ , it looks like:

$$[0, 0, \text{driver}[1]_x, \text{driver}[1]_y, \dots, 0, 0].$$

Under this setup, in Step 5, the server just need to evaluate the summation of all drivers' ciphertexts and subtract the rider's ciphertext from it, which gives a ciphertext which can be decrypted to:

$$[\text{driver}[0]_x - \text{rider}_x, \text{driver}[0]_y - \text{rider}_y, \text{driver}[1]_x - \text{rider}_x, \text{driver}[1]_y - \text{rider}_y, \dots]$$

, and then apply square inplace to obtain the evaluation ciphertext **result**.

Therefore, computing **result'**[ $2i$ ]+**result'**[ $2i+1$ ], the distance between the rider and  $i$ -th driver can be obtained immediately after the decryption (gives **result'**). The rider then can find the closest driver easily.

**Note.** Because of the limitation of batching slots available, We can have at most  $N/2$  drivers.

## 2.3. Serialization and Communication Channel

I didn't write a pair of client-server applications, but used two shared streams to simulate the communication between client and server. It can be transformed to two client-side and server-side applications with minimum efforts by replacing the dummy streams used here with a real data transfer stream for TCP/UDP.

- **parm\_stream**: used by the server to pass the parameters setting to clients.
- **data\_stream**: used by servers and clients to exchange public keys, linearization keys, and ciphertexts.

To optimize the communication package size transferred, I took use of serialization of the SEAL objects. Public keys, relinearization keys, symmetric encrypted ciphertexts in SEAL can all be serialized to have better space efficiency. However, public-key encrypted ciphertexts cannot be serialized. Therefore it is preferred to use the symmetric encryption if possible. That's why in Step 3, the rider is using her **sk** rather than **pk** to do the encryption.

## 2.4. Some Notable Observations

Here are some interesting findings that I observed when playing with my application:

1. The documentation of SEAL claims that the library built with **clang++** can run significantly faster than the one built with **g++**. But my program linked with them respectively does not show much difference given a relatively long execution time ~15s.
2. The noise budget remaining in the result ciphertext does not change much from maximum number of drivers (96 bits) i.e.  $N/2$  drivers to 3 drivers (101 bits).
3. The noise budget remaining in the result ciphertext does not change at all with or without relinearization after the additions and square operation. I can't explain this well.
4. A smaller plaintext modulus gives more remaining noise budget but has no impact on the overall running time.
5. A smaller polynomial modulus degree i.e.  $N = 4096$  with maximum drivers configured gives a much faster running speed from 15s to 2.5s, while at the cost of a much smaller remaining noise budget from 96 bits to 8 bits.

## 3. Conclusion

In this project, I experienced how to use a FHE library to implement an application with FHE in practice. I better understood the knowledge in lectures and learned how to play with 3rd party libraries in C++. I implemented an application for oblivious riding with batching and serialization.

## References

- [1] “Lattigo v2.1.0.” Online: <http://github.com/ldsec/lattigo>. Dec. 2020. EPFL-LDS.
- [2] “Microsoft SEAL (release 3.6).” <https://github.com/Microsoft/SEAL>. Nov. 2020. Microsoft Research, Redmond, WA.